

Nondeterministic Strategies and their Refinement in Strategy Logic

Giuseppe De Giacomo¹, Bastien Maubert² and Aniello Murano²

¹Sapienza Università di Roma

²Università degli Studi di Napoli Federico II

degiamaco@diag.uniroma1.it, bastien.maubert@gmail.com, murano@na.infn.it

Abstract

Nondeterministic strategies are strategies (or protocols, or plans) that, given a history in a game, assign a set of possible actions, all of which are winning. An important problem is that of refining such strategies. For instance, given a nondeterministic strategy that allows only safe executions, refine it to, additionally, eventually reach a desired state of affairs. We show that such problems can be solved elegantly in the framework of Strategy Logic (SL), a very expressive logic to reason about strategic abilities. Specifically, we introduce a variant of SL with nondeterministic strategies and a strategy refinement operator. We show that model checking this logic can be done at no additional computational cost with respect to standard SL, and can be used to solve problems synthesis, synthesis of most permissive strategies, module checking, and more.

1 Introduction

This paper is about nondeterministic strategies (aka plans or protocols), i.e., strategies that associate to the current history a *set of alternative moves* (instead of one) all of which are “good” for the objective of the strategy.

Nondeterministic strategies have been studied in literature in several context. Possibly the most relevant area is Discrete Event Control where a central notion is that of *maximally permissive supervisor* (). This is supervisor that controls a plant, i.e., allows the plant to do only certain operations at each point in time. Note that this supervisor does not says exactly what to do to the plant (as a deterministic strategy) but in fact tries to leave as much freedom as possible to plant itself blocking only operations that are unsafe. In fact it is of interest to be *maximally permissive* wrt the plant. And indeed the central result of Discrete Control Theory is that such a maximally permissive supervisor, i.e., nondeterministic strategy always exists if the plant and the supervisor specifications are expressed as regular languages.

Another interesting case is that controllers that orchestrate several components to compose a desired global behavior (?; ?). One way of seen this is that the controller tries to maintain overtime a sort of simulation relation between the desired behavior expressed as a transition system and the cartesian product of the transitino systems of the components.

Nondeterministic strategies are of interest in several context. For example, in planning when the action precondition specification can be seen as a nonseteterministi strategywe specify a function that given the state of the domain returns a set of possible actions. Now if we consider the state as summary of the relavant part of the history we can see

Although not as common as standard deterministic strategies, they are quite common

2 Nondeterministic strategies

We start with some basic notations, then we recall classic concurrent game structures, nondeterministic strategies, and the notion of strategy refinement.

2.1 Notations

Let Σ be an alphabet. A *finite* (resp. *infinite*) *word* over Σ is an element of Σ^* (resp. Σ^ω). The *length* of a finite word $w = w_0w_1 \dots w_n$ is $|w| := n + 1$, and $\text{last}(w) := w_n$ is its last letter. Given a finite (resp. infinite) word w and $0 \leq i < |w|$ (resp. $i \in \mathbb{N}$), we let w_i be the letter at position i in w , $w_{\leq i}$ is the prefix of w that ends at position i and $w_{>i}$ is the suffix that starts at position i . We write $w \preceq w'$ if w is a prefix of w' , and $\text{pref}(w)$ is the set of finite prefixes of word w . Finally, the domain of a mapping f is written $\text{dom}(f)$.

2.2 Concurrent game structures

For convenience we fix for the rest of the paper AP, a finite non-empty set of *atomic propositions*, and Ag, a finite non-empty set of *agents* or *players*.

Definition 1. A *concurrent game structure* (or CGS) is a tuple $\mathcal{G} = (\text{Ac}, V, E, \ell, v_i)$ where

- Ac is a finite non-empty set of *actions*,
- V is a finite non-empty set of *positions*,
- $E : V \times \text{Ac}^{\text{Ag}} \rightarrow V$ is a *transition function*,
- $\ell : V \rightarrow 2^{\text{AP}}$ is a *labelling function*, and
- $v_i \in V$ is an *initial position*.

In a position $v \in V$, where atomic propositions $\ell(v)$ hold, each player a chooses an action $c_a \in \text{Ac}$, and the game proceeds to position $E(v, c)$, where $c \in \text{Ac}^{\text{Ag}}$ stands for the *joint action* $(c_a)_{a \in \text{Ag}}$. Given a joint action $c = (c_a)_{a \in \text{Ag}}$

and $a \in \text{Ag}$, we let c_a denote c_a . A *finite* (resp. *infinite*) *play* is a finite (resp. infinite) word $\rho = v_0 \dots v_n$ (resp. $\pi = v_0 v_1 \dots$) such that $v_0 = v_i$ and for every i such that $0 \leq i < |\rho| - 1$ (resp. $i \geq 0$), there exists a joint action c such that $E(v_i, c) = v_{i+1}$. Given two finite plays ρ and ρ' , we say that ρ' is a *continuation* of ρ if $\rho' \in \rho \cdot V^*$, and we write $\text{Cont}(\rho)$ for the set of continuations of ρ .

[say that CGS also capture turn-based - Bastien]

2.3 Strategy refinement

Given a CGS \mathcal{G} , a *nondeterministic strategy*, or strategy for short, for a player is a function $\sigma : \text{Cont}(v_i) \rightarrow 2^{\text{Ac}} \setminus \emptyset$ that maps each finite play in \mathcal{G} to a nonempty finite set of actions that the player may choose from after this finite play. A strategy σ is *deterministic* if for every finite play ρ , $\sigma(\rho)$ is a singleton. We let Str denote the set of all (nondeterministic) strategies, and $\text{Str}^d \subset \text{Str}$ the set of deterministic ones (note that these sets depend on the CGS under consideration).

Formulas of our logic SL^\prec will be evaluated at the end of a finite play ρ (which can be simply the initial position of the game), and since SL^\prec contains only *future-time* temporal operators, the only relevant part of a strategy σ when evaluating a formula after finite play ρ is its definition on continuations of ρ . We thus define the *restriction* of σ to ρ as the restriction of σ to $\rho \cdot V^+$, that we write $\sigma|_\rho : \text{Cont}(\rho) \rightarrow 2^{\text{Ac}} \setminus \emptyset$. We will then say that a strategy σ *refines* another strategy σ' after a finite play ρ if the first one is more restrictive than the second one on continuations of ρ . More formally:

Definition 2. Strategy σ *refines* strategy σ' after finite play ρ , written $\sigma \preceq_\rho \sigma'$, if for every $\rho' \in \text{Cont}(\rho)$, $\sigma|_\rho(\rho') \subseteq \sigma'|_\rho(\rho')$. We simply say that σ *refines* σ' if it refines it after the initial position v_i , and in that case we write $\sigma \preceq \sigma'$.

[maximal permissive strategies: look at literature - Bastien]

3 Strategy Logic with refinement

In this section we introduce SL^\prec , which extends SL with nondeterministic strategies, an *outcome quantifier* that quantifies over possible outcomes of a strategy profile, and more importantly, a refining operator that expresses that a strategy refines another. We first fix some basic notations.

3.1 Syntax

In addition to the sets of propositions AP and agents Ag , we now fix Var , a finite non-empty set of *variables*.

Definition 3. The syntax of SL^\prec is defined by the following grammar:

$$\begin{aligned} \varphi &:= p \mid \neg\varphi \mid \varphi \vee \varphi \mid \exists x \varphi \mid x \preceq y \mid (a, x)\varphi \mid \mathbf{E}\psi \\ \psi &:= \varphi \mid \neg\psi \mid \psi \vee \psi \mid \mathbf{X}\psi \mid \psi \mathbf{U} \psi \end{aligned}$$

where $p \in \text{AP}$, $x, y \in \text{Var}$ and $a \in \text{Ag}$.

Formulas of type φ are called *state formulas*, those of type ψ are called *path formulas*, and SL^\prec consists of all state formulas.

Temporal operators, \mathbf{X} (next) and \mathbf{U} (until), have the usual meaning. The *refinement operator* expresses that the

strategy denoted by a variable x is more restrictive than another one, or that it allows less behaviours: $x \preceq y$ reads as “strategy x refines strategy y ”. The *strategy quantifier* $\exists x$ has its usual meaning, except that it now quantifies on *nondeterministic* strategies: $\exists x \varphi$ reads as “there exists a nondeterministic strategy x such that φ holds”, where x is a strategy variable. As usual, the *binding operator* (a, x) assigns a strategy to an agent, and $(a, x)\varphi$ reads as “when agent a plays strategy x , φ holds”. Finally, the *outcome quantifier* \mathbf{E} quantifies on outcomes of strategies currently in use: $\mathbf{E}\psi$ reads as “ ψ holds in some outcome of the strategies currently used by the players”.

We use usual abbreviations $\top := p \vee \neg p$, $\perp := \neg\top$, $\varphi \rightarrow \varphi' := \neg\varphi \vee \varphi'$, $\varphi \leftrightarrow \varphi' := \varphi \rightarrow \varphi' \wedge \varphi' \rightarrow \varphi$, $\mathbf{F}\varphi := \top \mathbf{U} \varphi$, $\mathbf{G}\varphi := \neg \mathbf{F} \neg \varphi$ and $\forall x \varphi := \neg \exists x \neg \varphi$.

For every formula $\varphi \in \text{SL}^\prec$, we let $\text{free}(\varphi)$ be the set of variables that appear free in φ , i.e., that appear out of the scope of a strategy quantifier. A formula φ is a *sentence* if $\text{free}(\varphi)$ is empty. Finally, we let the *size* $|\varphi|$ of a formula φ be the number of symbols in φ .

3.2 Semantics

SL^\prec formulas are interpreted in a CGS, and the semantics makes use of the following additional notions.

An *assignment* $\chi : \text{Ag} \cup \text{Var} \rightarrow \text{Str}$ is a partial function that assigns a strategy to each player and strategy variable in its domain. For an assignment χ , player a and strategy σ , $\chi[a \mapsto \sigma]$ is the assignment of domain $\text{dom}(\chi) \cup \{a\}$ that maps a to σ and is equal to χ on the rest of its domain, and $\chi[x \mapsto \sigma]$ is defined similarly, where x is a variable. An assignment is *variable-complete* for a formula $\varphi \in \text{SL}^\prec$ if its domain contains all free variables of φ .

For an assignment χ and a finite play ρ , we let $\text{Out}(\chi, \rho)$ be the set of infinite plays that start with ρ and are then extended by letting players follow the strategies assigned by χ . Formally, $\text{Out}(\chi, \rho)$ is the set of plays of the form $\rho \cdot v_1 v_2 \dots$ such that for all $i \geq 0$, there exists c such that for all $a \in \text{dom}(\chi) \cap \text{Ag}$, $c_a \in \chi(a)(\rho \cdot v_1 \dots v_i)$ and $v_{i+1} = E(v_i, c)$, with $v_0 = \text{last}(\rho)$.

Definition 4. The semantics of a state formula is defined on a CGS \mathcal{G} , an assignment χ that is variable-complete for φ , and a finite play ρ . For a path formula ψ , the finite play is replaced with an infinite play π and an index $i \in \mathbb{N}$. The

definition by mutual induction is as follows:

$\mathcal{G}, \chi, \rho \models p$	if $p \in \ell(\text{last}(\rho))$
$\mathcal{G}, \chi, \rho \models \neg \varphi$	if $\mathcal{G}, \chi, \rho \not\models \varphi$
$\mathcal{G}, \chi, \rho \models \varphi \vee \varphi'$	if $\mathcal{G}, \chi, \rho \models \varphi$ or $\mathcal{G}, \chi, \rho \models \varphi'$
$\mathcal{G}, \chi, \rho \models \exists x \varphi$	if $\exists \sigma \in \text{Str}$ s.t. $\mathcal{G}, \chi[x \mapsto \sigma], \rho \models \varphi$
$\mathcal{G}, \chi, \rho \models x \preceq y$	if $\chi(x)$ refines $\chi(y)$ after ρ
$\mathcal{G}, \chi, \rho \models (a, x) \varphi$	if $\mathcal{G}, \chi[a \mapsto \chi(x)], \rho \models \varphi$
$\mathcal{G}, \chi, \rho \models \mathbf{E}\psi$	if $\exists \pi \in \text{Out}(\chi, \rho)$ s.t. $\mathcal{G}, \chi, \pi, \rho - 1 \models \psi$
$\mathcal{G}, \chi, \pi, i \models \varphi$	if $\mathcal{G}, \chi, \pi_{\leq i} \models \varphi$
$\mathcal{G}, \chi, \pi, i \models \neg \psi$	if $\mathcal{G}, \chi, \pi, i \not\models \psi$
$\mathcal{G}, \chi, \pi, i \models \psi \vee \psi'$	if $\mathcal{G}, \chi, \pi, i \models \psi$ or $\mathcal{G}, \chi, \pi, i \models \psi'$
$\mathcal{G}, \chi, \pi, i \models \mathbf{X}\psi$	if $\mathcal{G}, \chi, \pi, i + 1 \models \psi$
$\mathcal{G}, \chi, \pi, i \models \psi \mathbf{U} \psi'$	if $\exists j \geq i$ s.t. $\mathcal{G}, \chi, \pi, j \models \psi'$ and, $\forall k$ s.t. $i \leq k < j$, $\mathcal{G}, \chi, \pi, k \models \psi$

We give some examples of useful notions that can be expressed in this logic.

Example 1 (Strategy equality). First, it is easy to see that a strategy σ refines a strategy σ' if $\sigma \preceq \sigma'$ and $\sigma' \preceq \sigma$. We thus define the abbreviation

$$x = y := x \preceq y \wedge y \preceq x$$

We thus have that $\mathcal{G}, \chi, \rho \models x = y$ if, and only if, $\chi(x)_{|\rho} = \chi(y)_{|\rho}$. And in particular, $\mathcal{G}, \chi, v_i \models x = y$ if, and only if, $\chi(x) = \chi(y)$. We also let $x \neq y := \neg(x = y)$ and $x \prec y := x \preceq y \wedge x \neq y$.

Example 2 (Deterministic strategies). We can also express that a strategy, or its refinement to continuations of the current finite play, is deterministic, with the following formula:

$$\text{Det}(x) := \forall y \ y \preceq x \rightarrow x \preceq y$$

Example 3 (Maximal permissive strategies). Given a formula $\varphi(x)$ we can express that a strategy x is maximally permissive with regards to $\varphi(x)$. Define formula $\text{MaxPerm}(x, \varphi)$ as follows:

$$\text{MaxPerm}(x, \varphi) := \varphi(x) \wedge (\forall y \ x \prec y \rightarrow \neg \varphi(y))$$

For instance, if we have two antagonistic players a and b , and a tries to ensure the safety property $\mathbf{G}p$, we can let $\varphi(x) = \forall z(a, x)(b, z)\mathbf{G}p$, and it then holds that $\mathcal{G}, \chi, v_i \models \text{MaxPerm}(x, \varphi)$ if, and only if, $\chi(x)$ is a maximally permissive winning strategy for a .

We now turn to establishing that the model-checking problem for SL^\prec is decidable. To do so we extend the classic approach, which is to reduce to QCTL^* , the extension of CTL^* with (second-order monadic) quantification on atomic propositions. This logic is equivalent to MSO on infinite trees (Laroussinie and Markey 2014), and it is easy to express that a strategy (or the atomic propositions that code for it) refines another one.

4 Synthesis

- Synthesis
- Max permissive
- plan B

5 Module checking

6 Refining Nash equilibria

7 Model checking SL^\prec

We first recall briefly the syntax and semantics of QCTL^* , to which we will reduce SL^\prec .

Definition 5. The syntax of QCTL^* is defined by the following grammar:

$$\begin{aligned} \varphi &:= p \mid \neg \varphi \mid \varphi \vee \varphi \mid \mathbf{E}\psi \mid \exists p \varphi \\ \psi &:= \varphi \mid \neg \psi \mid \psi \vee \psi \mid \mathbf{X}\psi \mid \psi \mathbf{U} \psi \end{aligned}$$

where $p \in \text{AP}$.

Again, formulas of type φ are called *state formulas*, those of type ψ are called *path formulas*, and QCTL^* consists of all the state formulas defined by the grammar, and we use standard abbreviation $\mathbf{A}\psi := \neg \mathbf{E} \neg \psi$.

The models of QCTL^* are classic Kripke structures:

Definition 6. A *Kripke structure*, or KS, over AP is a tuple $\mathcal{S} = (S, R, \ell, s_i)$ where

- S is a set of *states*,
- $R \subseteq S \times S$ is a left-total¹ *transition relation*,
- $\ell : S \rightarrow 2^{\text{AP}}$ is a *labelling function* and
- $s_i \in S$ is an *initial state*.

A *path* in \mathcal{S} is an infinite sequence of states $\lambda = s_0 s_1 \dots$ such that for all $i \in \mathbb{N}$, $(s_i, s_{i+1}) \in R$. A *finite path* is a finite non-empty prefix of a path. Similar to continuations of finite plays, given a finite path λ we write $\text{Cont}(\lambda)$ for the set of finite paths that start with λ . We may write $s \in \mathcal{S}$ for $s \in S$, and we define the *size* $|\mathcal{S}|$ of a KS $\mathcal{S} = (S, R, s_i, \ell)$ as its number of states: $|\mathcal{S}| := |S|$.

Since we will interpret QCTL^* on unfoldings of KS, we now define infinite trees.

Trees. Let X be a finite set of *directions* (typically a set of states). An *X-tree* τ is a nonempty set of words $\tau \subseteq X^+$ such that (1) there exists $r \in X$, called the *root* of τ , such that each $u \in \tau$ starts with r ($r \preceq u$); (2) if $u \cdot x \in \tau$ and $u \cdot x \neq r$, then $u \in \tau$; (3) if $u \in \tau$ then there exists $x \in X$ such that $u \cdot x \in \tau$.

The elements of a tree τ are called *nodes*. If $u \cdot x \in \tau$, we say that $u \cdot x$ is a *child* of u . An *X-tree* τ is *complete* if for every $u \in \tau$ and $x \in X$, $u \cdot x \in \tau$. A *path* in τ is an infinite sequence of nodes $\lambda = u_0 u_1 \dots$ such that for all $i \in \mathbb{N}$, u_{i+1} is a child of u_i , and $\text{Paths}(u)$ is the set of paths that start in node u .

An *AP-labelled X-tree*, or (AP, X) -tree for short, is a pair $t = (\tau, \ell)$, where τ is an *X-tree* called the *domain* of t and $\ell : \tau \rightarrow 2^{\text{AP}}$ is a *labelling*, which maps each node to the set of propositions that hold there. For $p \in \text{AP}$, a *p-labelling* for a tree is a mapping $\ell_p : \tau \rightarrow \{0, 1\}$ that indicates in which nodes p holds, and for a labelled tree $t = (\tau, \ell)$, the *p-labelling* of t is the *p-labelling* $u \mapsto 1$ if $p \in \ell(u)$, 0 otherwise. The composition of a labelled tree $t = (\tau, \ell)$ with a *p-labelling* ℓ_p for τ is defined as $t \otimes \ell_p := (\tau, \ell')$, where

¹i.e., for all $s \in S$, there exists s' such that $(s, s') \in R$.

$\ell'(u) = \ell(u) \cup \{p\}$ if $\ell_p(u) = 1$, and $\ell(u) \setminus \{p\}$ otherwise. A p -labelling for a labelled tree $t = (\tau, \ell)$ is a p -labelling for its domain τ . A *pointed labelled tree* is a pair (t, u) where u is a node of t .

Let $\mathcal{S} = (S, R, \ell, s_\iota)$ be a Kripke structure over AP. The *tree-unfolding* of \mathcal{S} is the (AP, S) -tree $t_{\mathcal{S}} := (\tau, \ell')$, where τ is the set of all finite paths that start in s_ι , and for every $u \in \tau$, $\ell'(u) := \ell(\text{last}(u))$.

Definition 7. We define by induction the satisfaction relation \models of QCTL*. Let $t = (\tau, \ell)$ be an AP-labelled tree, u a node and λ a path in τ :

$$\begin{aligned} t, u \models p & \quad \text{if} \quad p \in \ell(u) \\ t, u \models \neg\varphi & \quad \text{if} \quad t, u \not\models \varphi \\ t, u \models \varphi \vee \varphi' & \quad \text{if} \quad t, u \models \varphi \text{ or } t, u \models \varphi' \\ t, u \models \mathbf{E}\psi & \quad \text{if} \quad \exists \lambda \in \text{Paths}(u) \text{ s.t. } t, \lambda \models \psi \\ t, u \models \exists p \varphi & \quad \text{if} \quad \exists \ell_p \text{ a } p\text{-labelling for } t \text{ s.t.} \\ & \quad t \otimes \ell_p, u \models \varphi \\ t, \lambda \models \varphi & \quad \text{if} \quad t, \lambda_0 \models \varphi \\ t, \lambda \models \neg\psi & \quad \text{if} \quad t, \lambda \not\models \psi \\ t, \lambda \models \psi \vee \psi' & \quad \text{if} \quad t, \lambda \models \psi \text{ or } t, \lambda \models \psi' \\ t, \lambda \models \mathbf{X}\psi & \quad \text{if} \quad t, \lambda_{\geq 1} \models \psi \\ t, \lambda \models \psi \mathbf{U} \psi' & \quad \text{if} \quad \exists i \geq 0 \text{ s.t. } t, \lambda_{\geq i} \models \psi' \text{ and} \\ & \quad \forall j \text{ s.t. } 0 \leq j < i, t, \lambda_{\geq j} \models \psi \end{aligned}$$

We write $t \models \varphi$ for $t, r \models \varphi$, where r is the root of t . Given a KS \mathcal{S} and a QCTL* formula φ , we also write $\mathcal{S} \models \varphi$ if $t_{\mathcal{S}} \models \varphi$.

[define alternation depth - Bastien]

Theorem 1. The model-checking problem for QCTL* is $(k + 1)$ -EXPTIME-complete for formulas of alternation depth k .

7.1 Reduction to QCTL*

We use a variant of the reductions presented in (Laroussinie and Markey 2015; Fijalkow et al. 2018; Berthon et al. 2017; Maubert and Murano 2018; Bouyer et al. 2019), which transform instances of the model-checking problem for various strategic logics to (extensions of) QCTL*.

Let (\mathcal{G}, Φ) be an instance of the SL model-checking problem, and assume without loss of generality that each strategy variable is quantified at most once in Φ . We define an equivalent instance of the model-checking problem for QCTL*.

[say that the reduction preserves alternation depth - Bastien]

Define the KS $\mathcal{S}_{\mathcal{G}} := (S, R, s_\iota, \ell')$ where

- $S := \{s_v \mid v \in V\}$,
- $R := \{(s_v, s_{v'}) \mid \exists c \in \text{Ac}^{\text{Ag}} \text{ s.t. } E(v, c) = v'\} \subseteq S^2$,
- $s_\iota := s_{v_\iota}$, and
- $\ell'(s_v) := \ell(v) \cup \{p_v\} \subseteq \text{AP} \cup \text{AP}_v$.

For every finite play $\rho = v_0 \dots v_k$, define the node $u_\rho := s_{v_0} \dots s_{v_k}$ in $t_{\mathcal{S}_{\mathcal{G}}}$. Note that the mapping $\rho \mapsto u_\rho$ defines a bijection between the set of finite plays and the set of nodes in $t_{\mathcal{S}_{\mathcal{G}}}$.

Constructing the QCTL* formulas $(\varphi)_s^f$. We now describe how to transform an SL[<] formula φ and a partial function $f : \text{Ag} \rightarrow \text{Var}$ into a QCTL* formula $(\varphi)_s^f$ (that

will also depend on \mathcal{G}). Suppose that $\text{Ac} = \{c_1, \dots, c_l\}$, and define $(\varphi)_s^f$ and $(\psi)_p^f$ by mutual induction on state and path formulas. The base cases are as follows: $(p)_s^f := p$ and $(\varphi)_p^f := (\varphi)_s^f$. Boolean and temporal operators are simply obtained by distributing the translation: $(\neg\varphi)_s^f := \neg(\varphi)_s^f$, $(\neg\psi)_p^f := \neg(\psi)_p^f$, $(\varphi_1 \vee \varphi_2)_s^f := (\varphi_1)_s^f \vee (\varphi_2)_s^f$, $(\psi_1 \vee \psi_2)_p^f := (\psi_1)_p^f \vee (\psi_2)_p^f$, $(\mathbf{X}\psi)_p^f := \mathbf{X}(\psi)_p^f$ and $(\psi_1 \mathbf{U} \psi_2)_p^f := (\psi_1)_p^f \mathbf{U} (\psi_2)_p^f$.

We continue with the strategy quantifier:

$$(\exists x \varphi)_s^f := \exists p_{c_1}^x \dots \exists p_{c_l}^x. \varphi_{\text{str}}(x) \wedge (\varphi)_s^f$$

$$\text{where } \varphi_{\text{str}}(x) := \mathbf{AG} \bigvee_{c \in \text{Ac}} p_c^x$$

The intuition is that for each possible action $c \in \text{Ac}$, an existential quantification on the atomic proposition p_c^x “chooses” for each node u_ρ of the tree $t_{\mathcal{S}_{\mathcal{G}}}$ whether strategy x allows action c in ρ or not. $\varphi_{\text{str}}(x)$ checks that at least one action is allowed in each node, and thus that atomic propositions p_c^x indeed define a (nondeterministic) strategy.

For strategy refinement, the translation is as follows:

$$(x \preceq y)_s^f := \mathbf{AG} \bigwedge_{c \in \text{Ac}} p_c^x \rightarrow p_c^y.$$

Here are the remaining cases:

$$((a, x)\varphi)_s^f := (\varphi)_s^{f[a \mapsto x]} \quad \text{for } x \in \text{Var} \cup \{?\}$$

$$\text{and } (\mathbf{E}\psi)_s^f := \mathbf{E}(\psi_{\text{out}}^f \wedge (\psi)_p^f), \text{ where}$$

$$\begin{aligned} \psi_{\text{out}}^f &:= \mathbf{G} \bigvee_{v \in V} \left(p_v \wedge \right. \\ & \quad \left. \bigvee_{c \in \text{Ac}^{\text{Ag}}} \left(\bigwedge_{a \in \text{dom}(f)} p_{c_a}^{f(a)} \wedge \mathbf{X} p_{E(v, c)} \right) \right). \end{aligned}$$

ψ_{out}^f checks that each player a in the domain of f follows the strategy coded by the $p_c^{f(a)}$.

To prove the correctness of the translation we need some additional definitions. First, given a strategy σ and a strategy variable x we let $\ell_\sigma^x := \{\ell_{p_c^x} \mid c \in \text{Ac}\}$ be the family of p_c^x -labellings for tree $t_{\mathcal{S}_{\mathcal{G}}}$ defined as follows: for each finite play ρ and $c \in \text{Ac}$, we let $\ell_{p_c^x}(u_\rho) := 1$ if $c \in \sigma(\rho)$, 0 otherwise. For a labelled tree t with same domain as $t_{\mathcal{S}_{\mathcal{G}}}$ we write $t \otimes \ell_\sigma^x$ for $t \otimes \ell_{p_{c_1}^x} \otimes \dots \otimes \ell_{p_{c_l}^x}$.

Second, given an infinite play π and a point $i \in \mathbb{N}$, we let $\lambda_{\pi, i}$ be the infinite path in $t_{\mathcal{S}_{\mathcal{G}}}$ that starts in node $u_{\pi_{\leq i}}$ and is defined as $\lambda_{\pi, i} := u_{\pi_{\leq i}} u_{\pi_{\leq i+1}} u_{\pi_{\leq i+2}} \dots$.

Finally, we say that a partial function $f : \text{Ag} \rightarrow \text{Var}$ is *compatible* with an assignment χ if $\text{dom}(\chi) \cap \text{Ag} = \text{dom}(f)$ and for all $a \in \text{dom}(f)$, $\chi(a) = \chi(f(a))$.

Proposition 2. For every state subformula φ and path subformula ψ of Φ , finite play ρ , infinite play π , point $i \in \mathbb{N}$, for every assignment χ variable-complete for φ (resp. ψ) and partial function $f : \text{Ag} \rightarrow \text{Var}$ compatible with χ , assuming also that no x_i in $\text{dom}(\chi) \cap \text{Var} = \{x_1, \dots, x_k\}$ is quantified in φ or ψ , we have

$$\begin{aligned} \mathcal{G}, \chi, \rho \models \varphi & \quad \text{iff} \quad t_{\mathcal{S}_{\mathcal{G}}} \otimes \ell_{\chi(x_1)}^{x_1} \otimes \dots \otimes \ell_{\chi(x_k)}^{x_k}, u_\rho \models (\varphi)_s^f \\ \mathcal{G}, \chi, \pi, i \models \psi & \quad \text{iff} \quad t_{\mathcal{S}_{\mathcal{G}}} \otimes \ell_{\chi(x_1)}^{x_1} \otimes \dots \otimes \ell_{\chi(x_k)}^{x_k}, \lambda_{\pi, i} \models (\psi)_p^f \end{aligned}$$

In addition, \mathcal{S}_G is of size linear in $|\mathcal{G}|$, and $(\varphi)_s^f$ and $(\psi)_p^f$ are of size linear in $|\mathcal{G}|^2 + |\varphi|$.

Proof. The proof is by induction on φ . We detail the case for binding, strategy quantification, strategy refinement and outcome quantification, the others follow simply by definition of \mathcal{S}_G for atomic propositions and induction hypothesis for remaining cases.

For $\varphi = x \preceq y$, assume that $\mathcal{G}, \chi, \rho \models x \preceq y$. First, observe that since χ is variable-complete for φ , x and y are in $\text{dom}(\chi)$. Now we have that $\chi(x)_{|\rho}(\rho') \subseteq \chi(y)_{|\rho}(\rho')$ for every $\rho' \in \text{Cont}(\rho)$. By definition of $\ell_{\chi(x)}^x = \{\ell_{p_c^x} \mid c \in \text{Ac}\}$ and $\ell_{\chi(y)}^y = \{\ell_{p_c^y} \mid c \in \text{Ac}\}$, it follows that for each $c \in \text{Ac}$ and $\rho' \in \text{Cont}(\rho)$, if $\ell_{p_c^x}(\rho') = 1$, then $\ell_{p_c^y}(\rho') = 1$, and thus

$$t_{\mathcal{S}_G} \otimes \ell_{\chi(x)}^x \otimes \ell_{\chi(y)}^y \models \mathbf{AG} \bigwedge_{c \in \text{Ac}} p_c^x \rightarrow p_c^y$$

Because the labellings $\ell_{\chi(x)}^x$ touch distinct sets of atomic propositions for each variable x in $\text{dom}(\chi) \cap \text{Var}$, we can conclude this direction.

For the other direction let $t = t_{\mathcal{S}_G} \otimes \ell_{\chi(x_1)}^{x_1} \otimes \dots \otimes \ell_{\chi(x_k)}^{x_k}$ and assume that

$$t, u_\rho \models \mathbf{AG} \bigwedge_{c \in \text{Ac}} p_c^x \rightarrow p_c^y.$$

This implies that for every $\rho' \in \text{Cont}(\rho)$,

$$t, u_{\rho'} \models \bigwedge_{c \in \text{Ac}} p_c^x \rightarrow p_c^y,$$

and thus $\chi(x)_{|\rho}$ refines $\chi(y)_{|\rho}$.

For $\varphi = (a, x)\varphi'$, we have $\mathcal{G}, \chi, \rho \models (a, x)\varphi'$ if and only if $\mathcal{G}, \chi[a \mapsto \chi(x)], \rho \models \varphi'$. The result follows by using the induction hypothesis with assignment $\chi[a \mapsto x]$ and function $f[a \mapsto x]$. This is possible because $f[a \mapsto x]$ is compatible with $\chi[a \mapsto x]$: indeed $\text{dom}(\chi[a \mapsto x]) \cap \text{Ag}$ is equal to $\text{dom}(\chi) \cap \text{Ag} \cup \{a\}$ which, by assumption, is equal to $\text{dom}(f) \cup \{a\} = \text{dom}(f[a \mapsto x])$. Also by assumption, for all $a' \in \text{dom}(f)$, $\chi(a') = \chi(f(a'))$, and by definition $\chi[a \mapsto \chi(x)](a) = \chi(x) = \chi(f[a \mapsto x](a))$.

For $\varphi = \exists x \varphi'$, assume first that $\mathcal{G}, \chi, \rho \models \exists x \varphi'$. There exists a nondeterministic strategy σ such that

$$\mathcal{G}, \chi[x \mapsto \sigma], \rho \models \varphi'.$$

Since f is compatible with χ , it is also compatible with assignment $\chi' = \chi[x \mapsto \sigma]$. By assumption, no variable in $\{x_1, \dots, x_k\}$ is quantified in φ , so that $x \neq x_i$ for all i , and thus $\chi'(x_i) = \chi(x_i)$ for all i ; and because no strategy variable is quantified twice in a same formula, x is not quantified in φ' , so that no variable in $\{x_1, \dots, x_k, x\}$ is quantified in φ' . By induction hypothesis

$$t_{\mathcal{S}_G} \otimes \ell_{\chi'(x_1)}^{x_1} \otimes \dots \otimes \ell_{\chi'(x_k)}^{x_k} \otimes \ell_{\chi'(x)}^x, u_\rho \models (\varphi')_s^f.$$

It follows that

$$t_{\mathcal{S}_G} \otimes \ell_{\chi(x_1)}^{x_1} \otimes \dots \otimes \ell_{\chi(x_k)}^{x_k}, u_\rho \models \exists \tilde{\sigma} p_{c_1}^x \dots \exists p_{c_l}^x \cdot \varphi_{\text{str}}(x) \wedge (\varphi')_s^f.$$

Finally, since $\chi'(x_i) = \chi(x_i)$ for all i , we conclude that

$$t_{\mathcal{S}_G} \otimes \ell_{\chi(x_1)}^{x_1} \otimes \dots \otimes \ell_{\chi(x_k)}^{x_k}, u_\rho \models (\exists x \varphi')_s^f.$$

For the other direction, assume that

$$t_{\mathcal{S}_G} \otimes \ell_{\chi(x_1)}^{x_1} \otimes \dots \otimes \ell_{\chi(x_k)}^{x_k}, u_\rho \models (\varphi)_s^f,$$

and recall that $(\varphi)_s^f = \exists \tilde{\sigma} p_{c_1}^x \dots \exists \tilde{\sigma} p_{c_l}^x \cdot \varphi_{\text{str}}(x) \wedge (\varphi')_s^f$. Write $t = t_{\mathcal{S}_G} \otimes \ell_{\chi(x_1)}^{x_1} \otimes \dots \otimes \ell_{\chi(x_k)}^{x_k}$. There exist $\ell_{p_c^x}$ -labellings such that

$$t \otimes \ell_{p_{c_1}^x} \otimes \dots \otimes \ell_{p_{c_l}^x} \models \varphi_{\text{str}}(x) \wedge (\varphi')_s^f.$$

By $\varphi_{\text{str}}(x)$, these labellings code for a strategy σ . Let $\chi' = \chi[x \mapsto \sigma]$. For all $1 \leq i \leq k$, by assumption $x \neq x_i$, and thus $\chi'(x_i) = \chi(x_i)$. The above can thus be rewritten

$$t_{\mathcal{S}_G} \otimes \ell_{\chi'(x_1)}^{x_1} \otimes \dots \otimes \ell_{\chi'(x_k)}^{x_k} \otimes \ell_{\chi'(x)}^x \models \varphi_{\text{str}}(x) \wedge (\varphi')_s^f.$$

By induction hypothesis we have $\mathcal{G}, \chi[x \mapsto \sigma], \rho \models \varphi'$, hence $\mathcal{G}, \chi, \rho \models \exists x \varphi'$.

For $\varphi = \exists^d x \psi$, the proof is similar, using $\varphi_{\text{str}}^{\text{det}}(x)$ instead of $\varphi_{\text{str}}(x)$.

For $\varphi = \mathbf{E}\psi$, assume first that $\mathcal{G}, \chi, \rho \models \mathbf{E}\psi$. There exists a play $\pi \in \text{Out}(\chi, \rho)$ s.t. $\mathcal{G}, \chi, \pi, |\rho| - 1 \models \psi$. By induction hypothesis,

$$t_{\mathcal{S}_G} \otimes \ell_{\chi(x_1)}^{x_1} \otimes \dots \otimes \ell_{\chi(x_k)}^{x_k}, \lambda_{\pi, |\rho| - 1} \models (\psi)_p^f.$$

Since π is an outcome of χ , each agent $a \in \text{dom}(\chi) \cap \text{Ag}$ follows strategy $\chi(a)$ in π . Because $\text{dom}(\chi) \cap \text{Ag} = \text{dom}(f)$ and for all $a \in \text{dom}(f)$, $\chi(a) = \chi(f(a))$, each agent $a \in \text{dom}(f)$ follows the strategy $\chi(f(a))$, which is coded by atoms $p_c^{f(a)}$ in the translation of Φ . Therefore $\lambda_{\pi, |\rho| - 1}$ also satisfies ψ_{out}^x , hence

$$t_{\mathcal{S}_G} \otimes \ell_{\chi(x_1)}^{x_1} \otimes \dots \otimes \ell_{\chi(x_k)}^{x_k}, \lambda_{\pi, |\rho| - 1} \models \psi_{\text{out}}^x \wedge (\psi)_p^f,$$

and we are done.

For the other direction, assume that

$$t_{\mathcal{S}_G} \otimes \ell_{\chi(x_1)}^{x_1} \otimes \dots \otimes \ell_{\chi(x_k)}^{x_k}, u_\rho \models \mathbf{E}(\psi_{\text{out}}^f \wedge (\psi)_p^f).$$

There exists a path λ in $t_{\mathcal{S}_G} \otimes \ell_{\chi(x_1)}^{x_1} \otimes \dots \otimes \ell_{\chi(x_k)}^{x_k}$ starting in node u_ρ that satisfies both ψ_{out}^f and $(\psi)_p^f$. By construction of \mathcal{S}_G there exists an infinite play π such that $\pi_{\leq |\rho| - 1} = \rho$ and $\lambda = \lambda_{\pi, |\rho| - 1}$. By induction hypothesis, $\mathcal{G}, \chi, \pi, |\rho| - 1 \models \psi$. Because $\lambda_{\pi, |\rho| - 1}$ satisfies ψ_{out}^f , $\text{dom}(\chi) \cap \text{Ag} = \text{dom}(f)$, and for all $a \in \text{dom}(f)$, $\chi(a) = \chi(f(a))$, it is also the case that $\pi \in \text{Out}(\chi, \rho)$, hence $\mathcal{G}, \chi, \rho \models \mathbf{E}\psi$. \square

Applying Proposition 2 to the sentence Φ , $\rho = v_i$, any assignment χ , and the empty function \emptyset , we get:

$$\mathcal{G} \models \Phi \quad \text{if and only if} \quad t_{\mathcal{S}_G} \models (\Phi)_s^f.$$

We now show how SL^\prec captures various a number of important problems.

References

- Berthon, R.; Maubert, B.; Murano, A.; Rubin, S.; and Vardi, M. Y. 2017. Strategy logic with imperfect information. In *LICS'17*, 1–12. IEEE.
- Bouyer, P.; Kupferman, O.; Markey, N.; Maubert, B.; Murano, A.; and Perelli, G. 2019. Reasoning about quality and fuzziness of strategic behaviours. In *Proceedings of the Twenty-Eighth International Joint Conference on Artificial Intelligence, IJCAI 2019, Macao, China, August 10-16, 2019*, 1588–1594.
- Fijalkow, N.; Maubert, B.; Murano, A.; and Rubin, S. 2018. Quantifying bounds in strategy logic. In *27th EACSL Annual Conference on Computer Science Logic, CSL 2018, September 4-7, 2018, Birmingham, UK*, 23:1–23:23.
- Laroussinie, F., and Markey, N. 2014. Quantified CTL: expressiveness and complexity. *LMCS* 10(4).
- Laroussinie, F., and Markey, N. 2015. Augmenting ATL with strategy contexts. *Inf. Comput.* 245:98–123.
- Maubert, B., and Murano, A. 2018. Reasoning about knowledge and strategies under hierarchical information. In *Principles of Knowledge Representation and Reasoning: Proceedings of the Sixteenth International Conference, KR 2018, Tempe, Arizona, 30 October - 2 November 2018*, 530–540.